

1. Introduction

In reinforcement learning it is common to let an agent interact with its environment for a **fixed amount of time** before resetting the environment and repeating the process in a series of episodes.

The task can either be to maximize the performance over

- ▶ that fixed period → **time-limited task**

$$G_{t:T}^{\gamma} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T = \sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1} \quad (1)$$

- ▶ an infinite period (unless environmental termination) while time limits are used during training to diversify experience → **time-unlimited task**

$$G_t^{\gamma} \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \quad (2)$$

2. Standard approach

In the **standard** benchmark domains (e.g. OpenAI Gym) or algorithms (e.g. OpenAI Baselines), this distinction is often overlooked.

- ▶ The state-value function is defined as:

$$v_{\pi}(s) \doteq \mathbb{E}_{\pi} [G_t^{\gamma} | S_t = s] \quad (3)$$

- ▶ The temporal-difference targets to approximate $v_{\pi}(s)$ are:

$$\begin{aligned} r & \text{ at termination (including time limits)} \\ r + \gamma \hat{v}_{\pi}(s) & \text{ otherwise} \end{aligned} \quad (4)$$

Issues:

- ▶ Terminations due to time limits are perceived by the agent as environmental terminations that **dynamically change** with the agent's policy.
- ▶ The unperceived remaining time leads to **state aliasing**.

3. Proposed approaches

Time-awareness for time-limited tasks:

Include a notion of the **remaining time in input**.

- ▶ The state-value function is defined as:

$$v_{\pi}(s, T - t) \doteq \mathbb{E}_{\pi} [G_{t:T}^{\gamma} | S_t = s] \quad (5)$$

- ▶ The temporal-difference targets for $v_{\pi}(s, T - t)$ are:

$$\begin{aligned} r & \text{ at termination (including time limits)} \\ r + \gamma \hat{v}_{\pi}(s', T - t - 1) & \text{ otherwise} \end{aligned} \quad (6)$$

Partial-episode bootstrapping for time-unlimited tasks:

Bootstrap when termination is due to time limits.

- ▶ The state-value function is defined as:


$$v_{\pi}(s) \doteq \mathbb{E}_{\pi} [G_{t:T}^{\gamma} + \gamma^{T-t} v_{\pi}(S_T) | S_t = s] \quad (7)$$

- ▶ The temporal-difference targets for $\hat{v}_{\pi}(s)$ are:


$$\begin{aligned} r & \text{ at } \textbf{environmental} \text{ termination} \\ r + \gamma \hat{v}_{\pi}(s') & \text{ otherwise } \textbf{(including time limits)} \end{aligned} \quad (8)$$

4. A motivational example


The Two-Goal Gridworld problem: Two rewarding terminal states (50 top-right and 20 bottom-left), a penalty of -1 for moving, a time limit $T = 3$.



Standard



Time-aware




Partial-ep. bootstrap



Policies and color-coded state-values. The **standard** agent optimizes for none of the two tasks. The agent with **time-awareness**, that learns to stay in place when there is not enough time to reach a goal, optimizes for the time-limited task. The agent with **partial-episode bootstrapping**, that aims for the highest reward, optimizes for the time-unlimited task.

5. Time-awareness

The Last Moment problem: The agent has to jump just before the time limit to maximize its score. Only a **time-aware** agent can learn the optimal policy for $T > 1$.




The Queue of Cars problem: The agent has to reach the exit before the time limit. The “safe” action moves the car with 50% probability, the “dangerous” action does it with 80% probability but with 10% chance of collision. A reward of 1 is provided at destination.

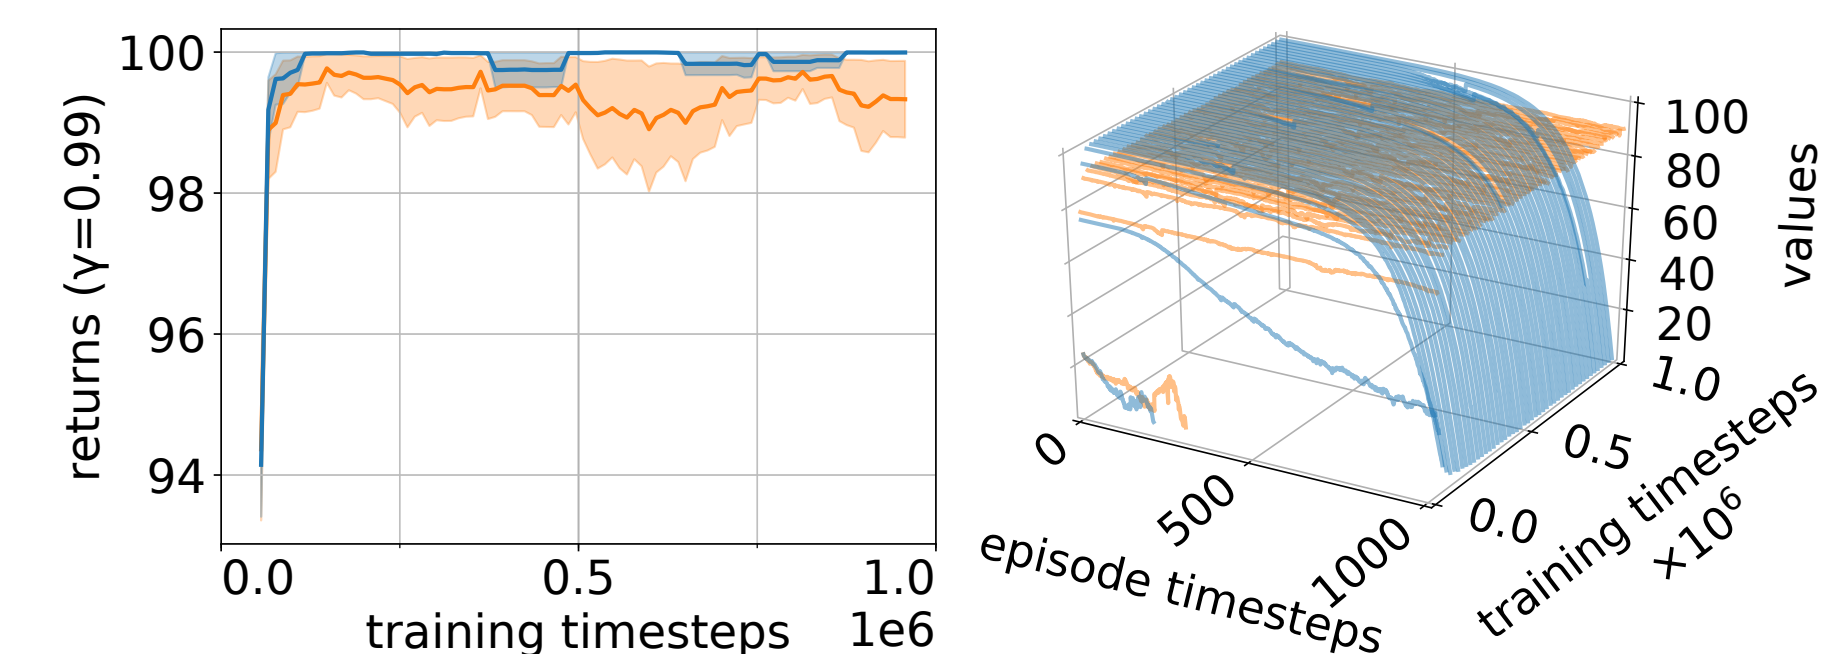


Decay ($\gamma = 0.99$)


Reacher-v1 $T = 50$



InvertedPendulum-v1 $T = 1000$




Hopper-v1 $T = 300$



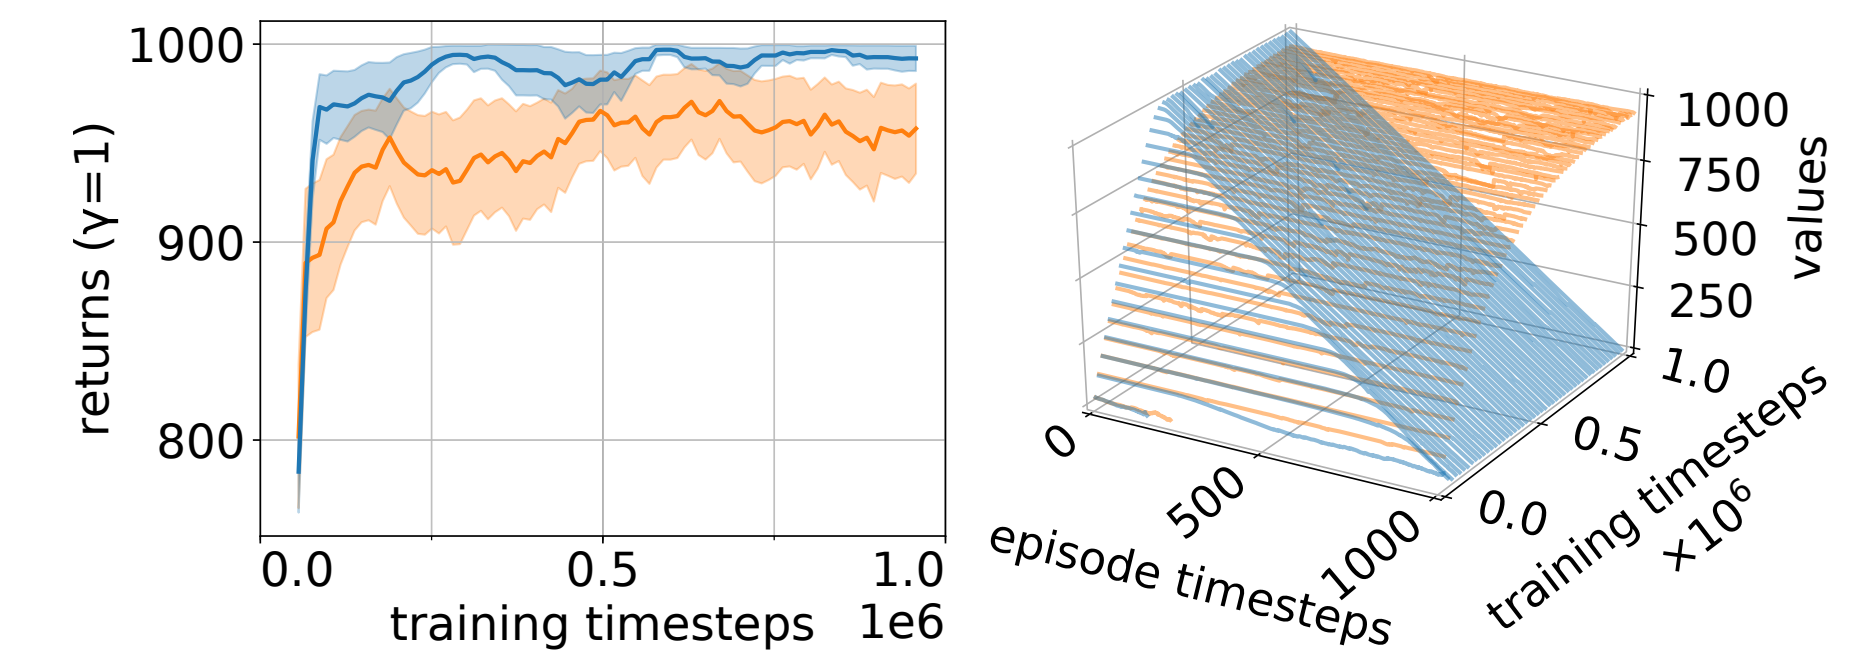
Time-aware PPO Standard PPO

No decay ($\gamma = 1$)


Reacher-v1 $T = 50$



InvertedPendulum-v1 $T = 1000$



Hopper-v1 $T = 300$




Time-aware PPO Standard PPO

Implementation details for the proposed agents:

- ▶ Two-Goal Gridworld problem: three Q-learning tables were individually learned, one for each time step.
- ▶ Other experiments: a scalar representing the normalized remaining time was concatenated to the observations.



6. Partial-episode bootstrapping

Hopper-v1: Trained with a small time limit ($T = 200$, i.e. 2–3 hops) and tested with a very large time limit ($T = 10^6$, i.e. more than 2 hours of rendered hopping). With **partial-episode bootstrapping**, PPO manages to reach the evaluation time limit multiple times.



The Infinite Cube Pusher task:

The agent has to push the cube onto the target with the ball. The target moves to a new position and a reward of 1 is given every time. Trained with a small time limit ($T = 50$, i.e. 1 push) and tested with a larger time limit ($T = 1000$). With **partial-episode bootstrapping**, PPO manages to reach 36 targets.



Implementation details for the proposed agent:

- ▶ The $GAE(\lambda)$ of PPO was modified to keep bootstrapping when termination was only due to time limits.

Information

Article and videos are available at: sites.google.com/view/time-limits-in-rl

Contact: f.pardo@imperial.ac.uk

Research presented in this poster has been supported by Dyson Technology Ltd.